

Cloud Computing

Course plan

Learning session 2-3

Core Services

Disclaimer

- Most Clouds completely abstract *real* infrastructure
- End-users deal with *services* on top of real infrastructure

“Core Services” lecture plan

- Compute
- Storage
- Network
- Databases
- Events and Messaging

Compute

Compute is ...

- infrastructure capable to run your workload/application
- CPU + memory

Compute

Compute could be based on:

- *raw* Virtual Machines (VM)
- Functions (Function as a Service)
- Containers
- Mixed/Hybrid

Compute

Virtual Machine

- Configurable CPU, RAM, size/type of storage, network bandwidth
- Control OS components and configuration
- Create snapshots/images

Virtual Machines / EC2 / Compute Engine

Compute

Functions

- Run **code** in response to event (http or from other services)
- No VMs to manage
- Faster scale-out
- (often) Pay only for running time (ms)

Azure Function / AWS Lambda / Cloud Function

Compute

Containers

- Run single **container** without orchestrator
- No VMs to manage

Azure Container Instance / AWS Fargate / Cloud Run

Compute

Kubernetes:

- Run **multiple containers on** (partially) cloud-managed **Kubernetes**
- Integrated with other cloud services
- (mostly) based on VMs

AKS / EKS / GKE

Compute

Application-optimized:

- **VM** optimized for concrete app
- Pre-installed applications/services
- Often includes other cloud services (monitoring, dns, and others)
- Cloud-provider offers marketplace

App Service / Lightsail or Elastic Beanstalk / App Engine

Compute

Batch:

- **VMs** optimized for high performance computing

Azure Batch / AWS Batch / Preemptible Virtual Machines

Compute

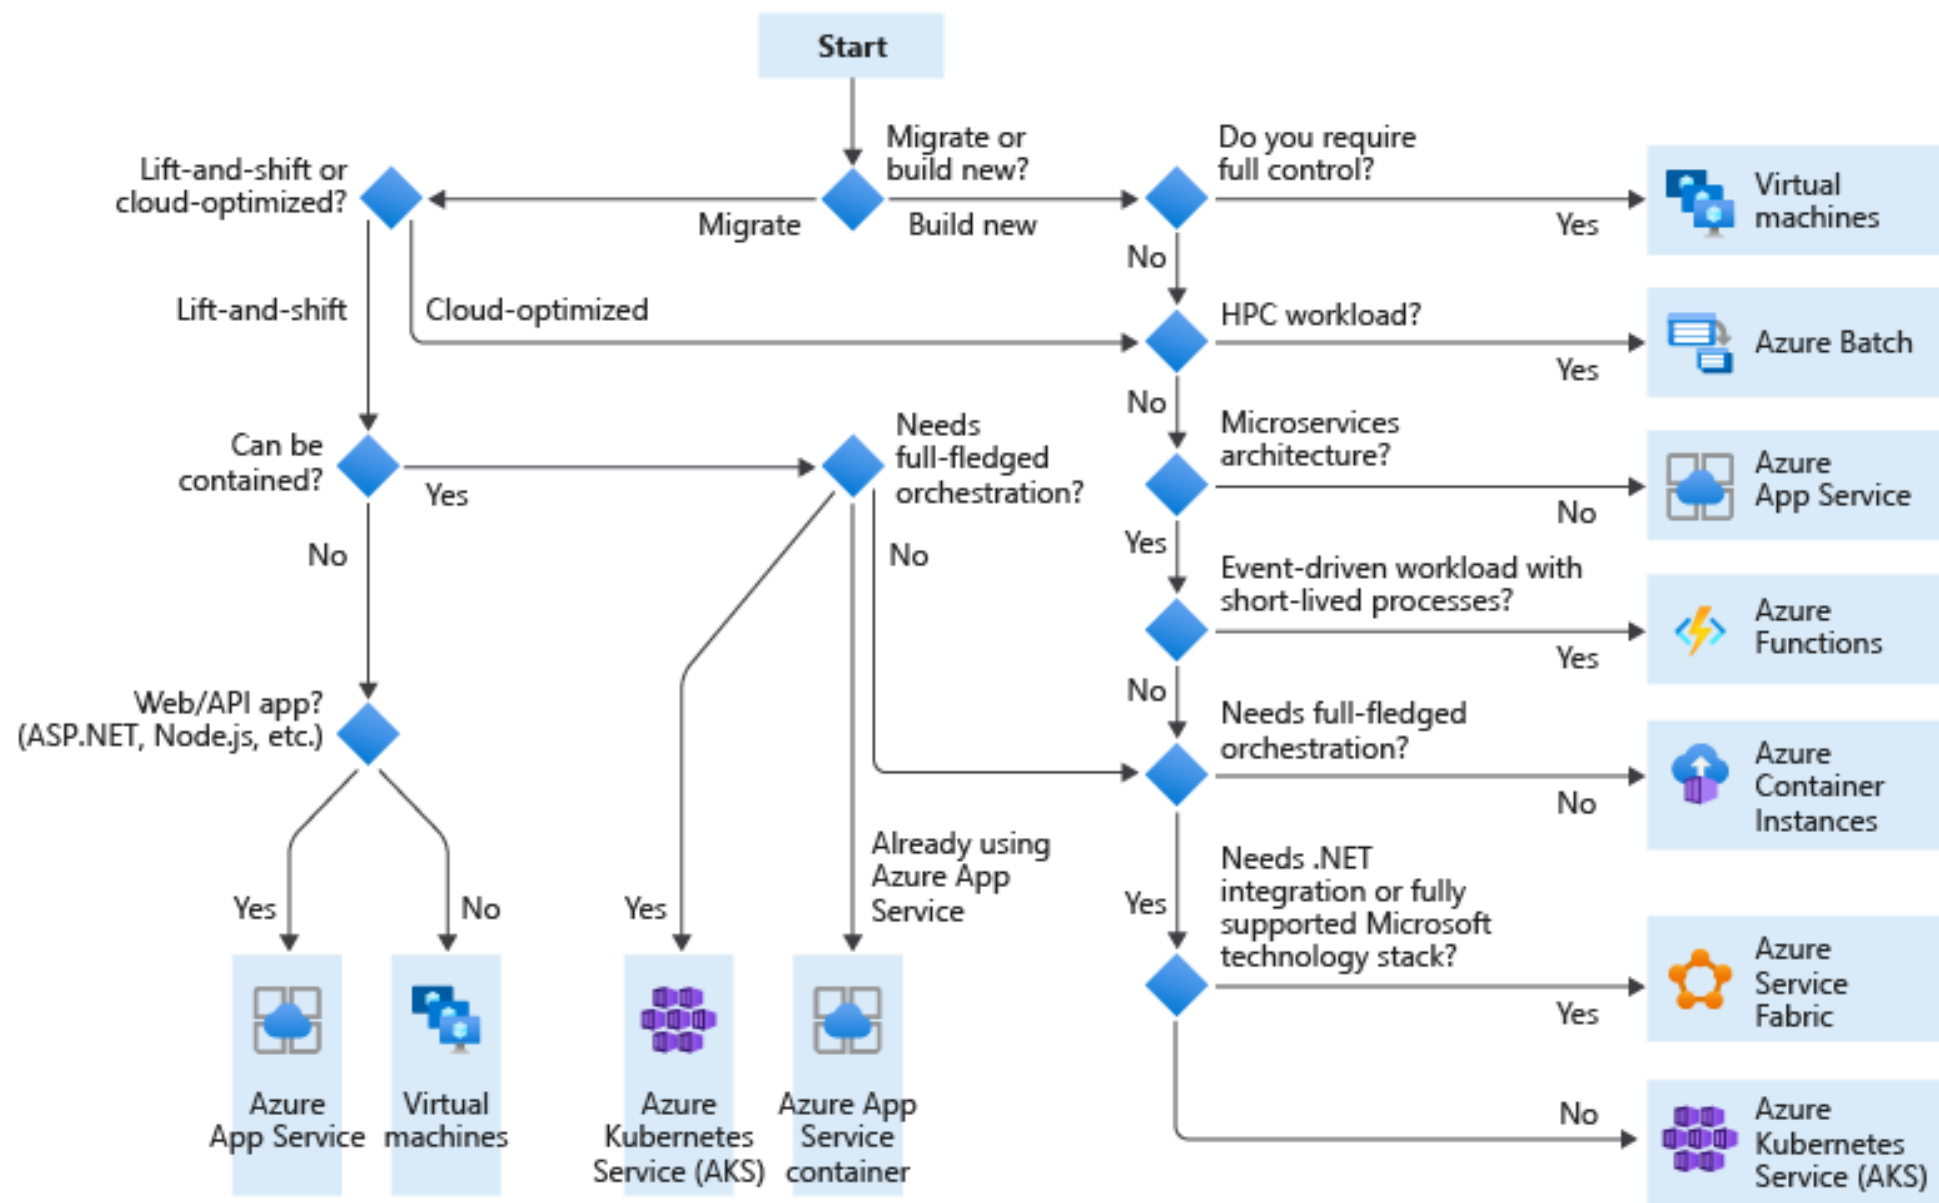
Other:

- [Service Fabric](#): distributed systems platform (azure's k8s)
- ???

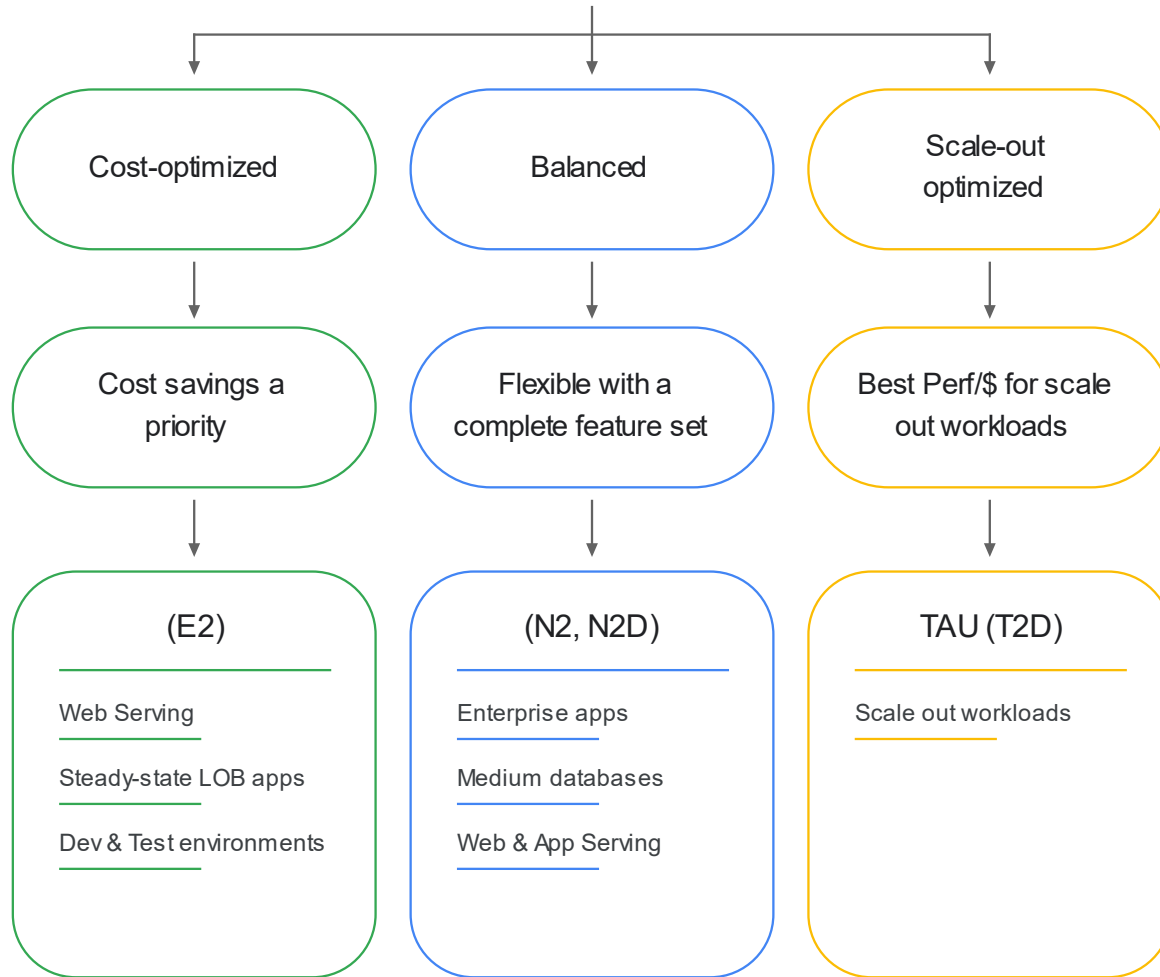
Compute

How to choose compute service?

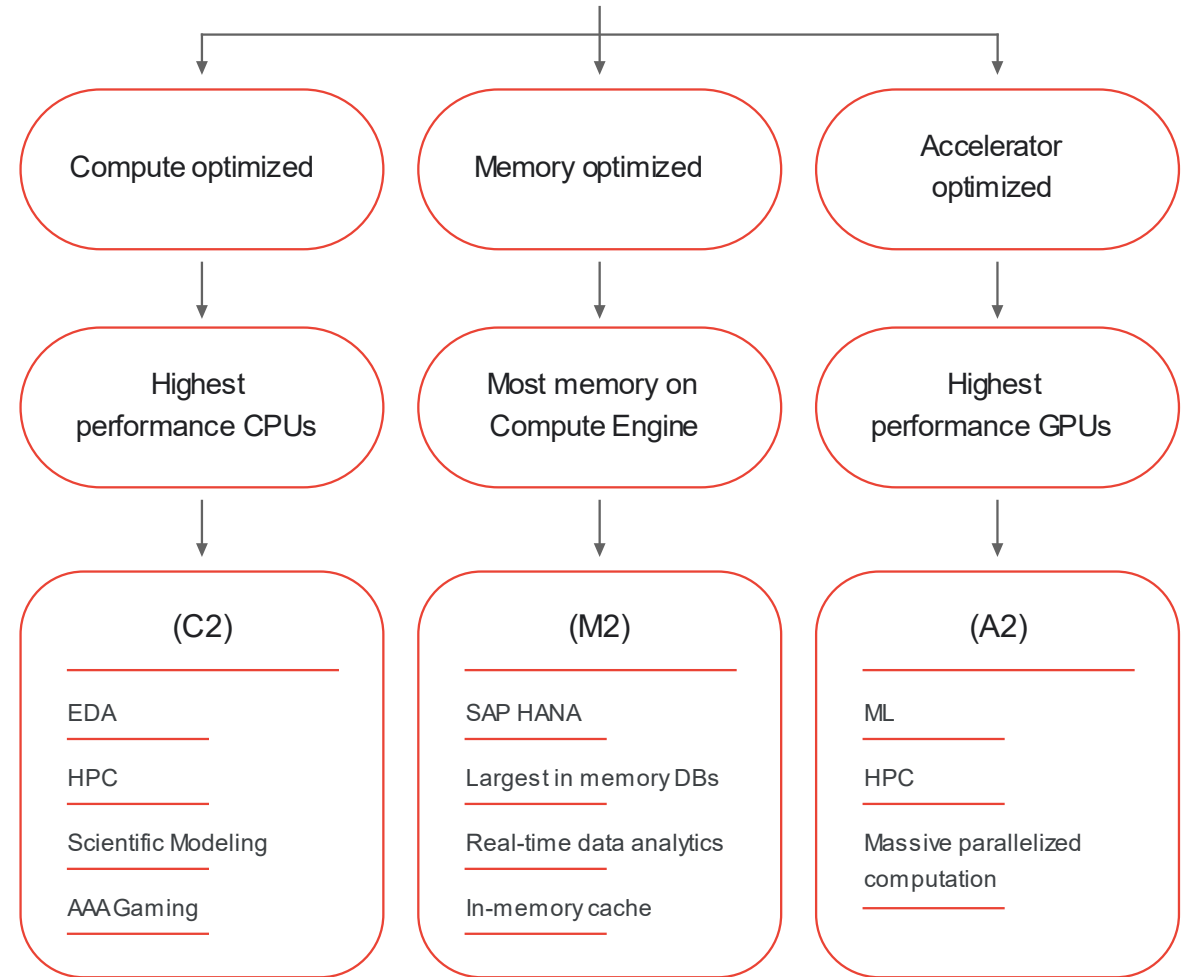
- [Azure](#)
- [AWS](#)
- [GCP compute](#)



General purpose



Workload-optimized



Compute

Demo

Compute Summary

- **Compute** is an infrastructure capable to run workload/application
- Main types:
 - VM
 - Function
 - Container

Storage

Storage is an infrastructure capable to store your data

Storage

Types of storage resources

- Block
- File system
- Object

Storage

Block Storage:

- Specific location on disks/memory
- Network access to raw block devices
- Each block could be referenced in different OS and partitioned independently

(Solid State Drive)




(Hard Disk Drive)



Storage

File Storage:


- Shared filesystem
- Folder structure
- File path and name



Storage

Object Storage:

- Flat-structure like key-value store
- Each object has associated metadata and Id
- Immutable objects: written or replaced
- Often accessed via (REST) API
- Often allows aggregation into buckets/containers



Storage

AWS offerings:

- Block Storage: **EBS** (Elastic Block Storage)
- File System: **EFS** (Elastic File System)
- Object Storage: **S3** (Simple Storage Service) with multiple replication types and access tiers ([Storage Classes](#)) including **Glacier**

Bonus: **Snow Family**

Storage

Azure offerings:

- Block Storage: [Azure Disk](#)
- File System: [Azure File](#) or [One Drive](#)
- Object Storage: [Storage Account](#), including [Blob](#), [Queue](#), [Table](#), [File Share](#) (File System)
 - Has multiple access tiers and replication modes

Storage

GCP offerings:

- Block Storage: [Persistent Disk](#)
- File System: [Filestore](#) or [Google Drive](#)
- Object Storage: [Cloud Storage](#) with multiple access tiers ([Storage Classes](#)) and redundancy options

Storage

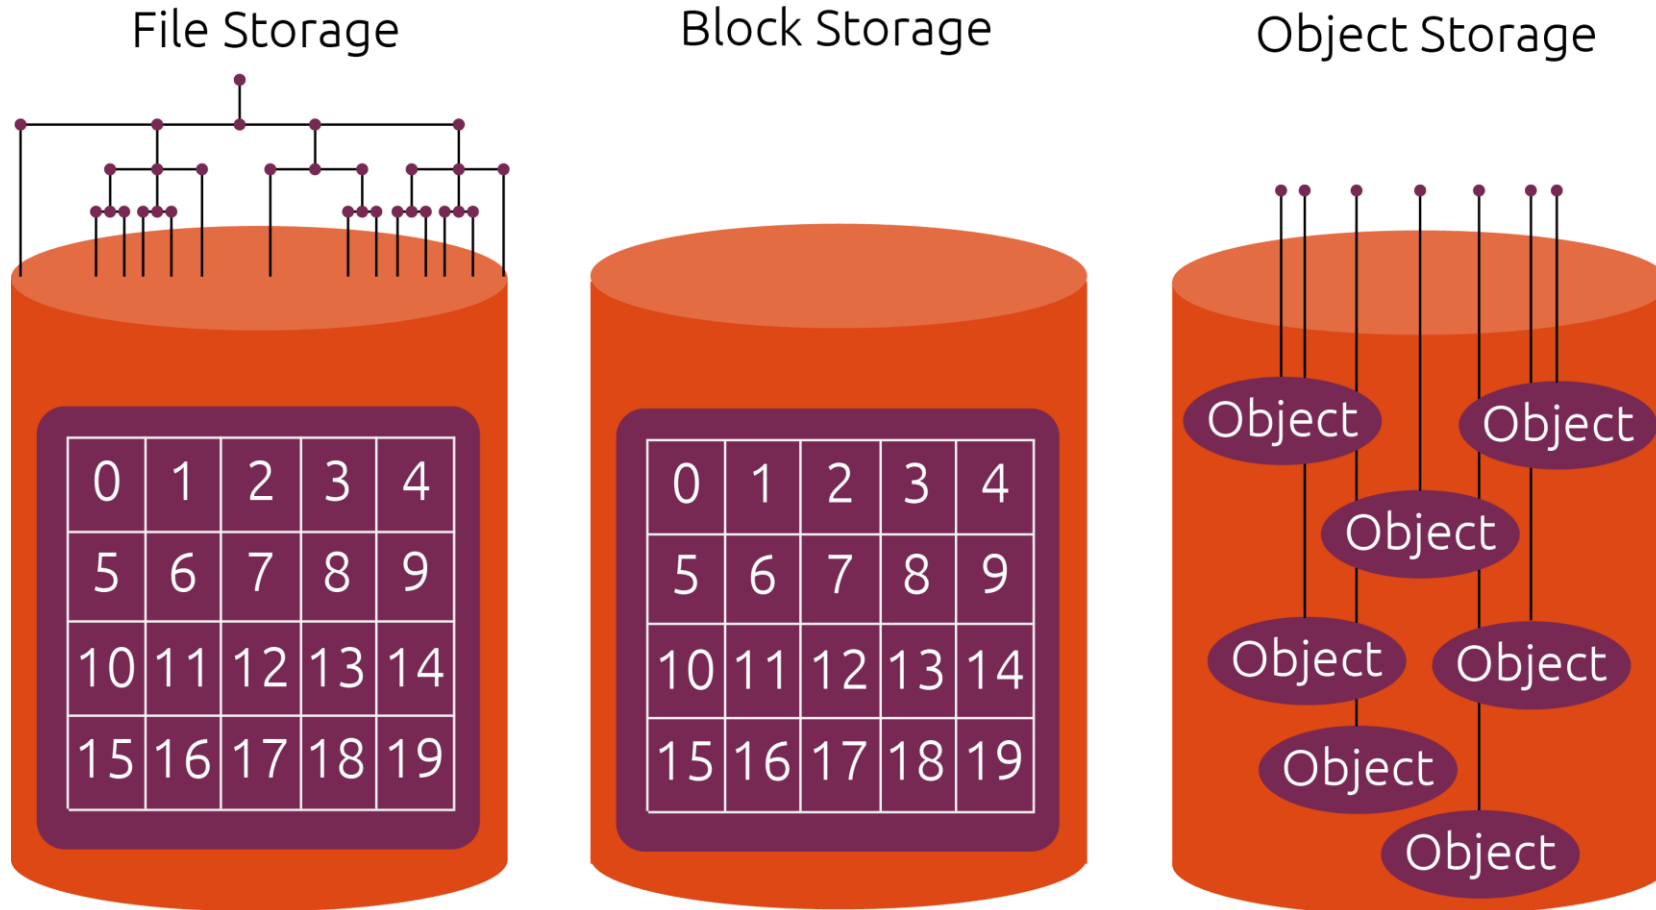
How to choose:

- Who uses the storage
- Usage patterns and performance: access frequency, size of data, size of files
- Pricing
- Access control

Storage

Demo

Storage Summary



Storage

Examples

- VM disk
- DB Backup
- A bunch of json files
- Data Lake
- Share files within a team

[Bonus] Storage



Network

- Cloud does not offer a physical network, but it provides networking services
- Networking services – are the ones that solve your **connectivity** and **isolation** problems
- [Software-defined network](#)

Network

Classless Inter-Domain Routing (**CIDR**)

- Used for IP-allocation and routing
- It divides the space using 1-bit steps
- Consists of two parts:
 - the most significant bit
 - network prefix

Network

CIDR notation	Bits	Range
10.0.0.0/8	00001010.00000000.00000000.00000000	10.0.0.0 – 10.255.255.255
10.1.0.0/16	00001010.00000001.00000000.00000000	10.1.0.0 – 10.1.255.255
10.2.0.0/16	00001010.00000010.00000000.00000000	10.2.0.0 – 10.2.255.255
10.3.0.0/31	00001010.00000011.00000000.00000000	10.3.0.0 – 10.3.0.1
10.3.0.2/32	00001010.00000011.00000000.00000010	10.3.0.2 – 10.3.0.2

CIDR [Calculator example](#)

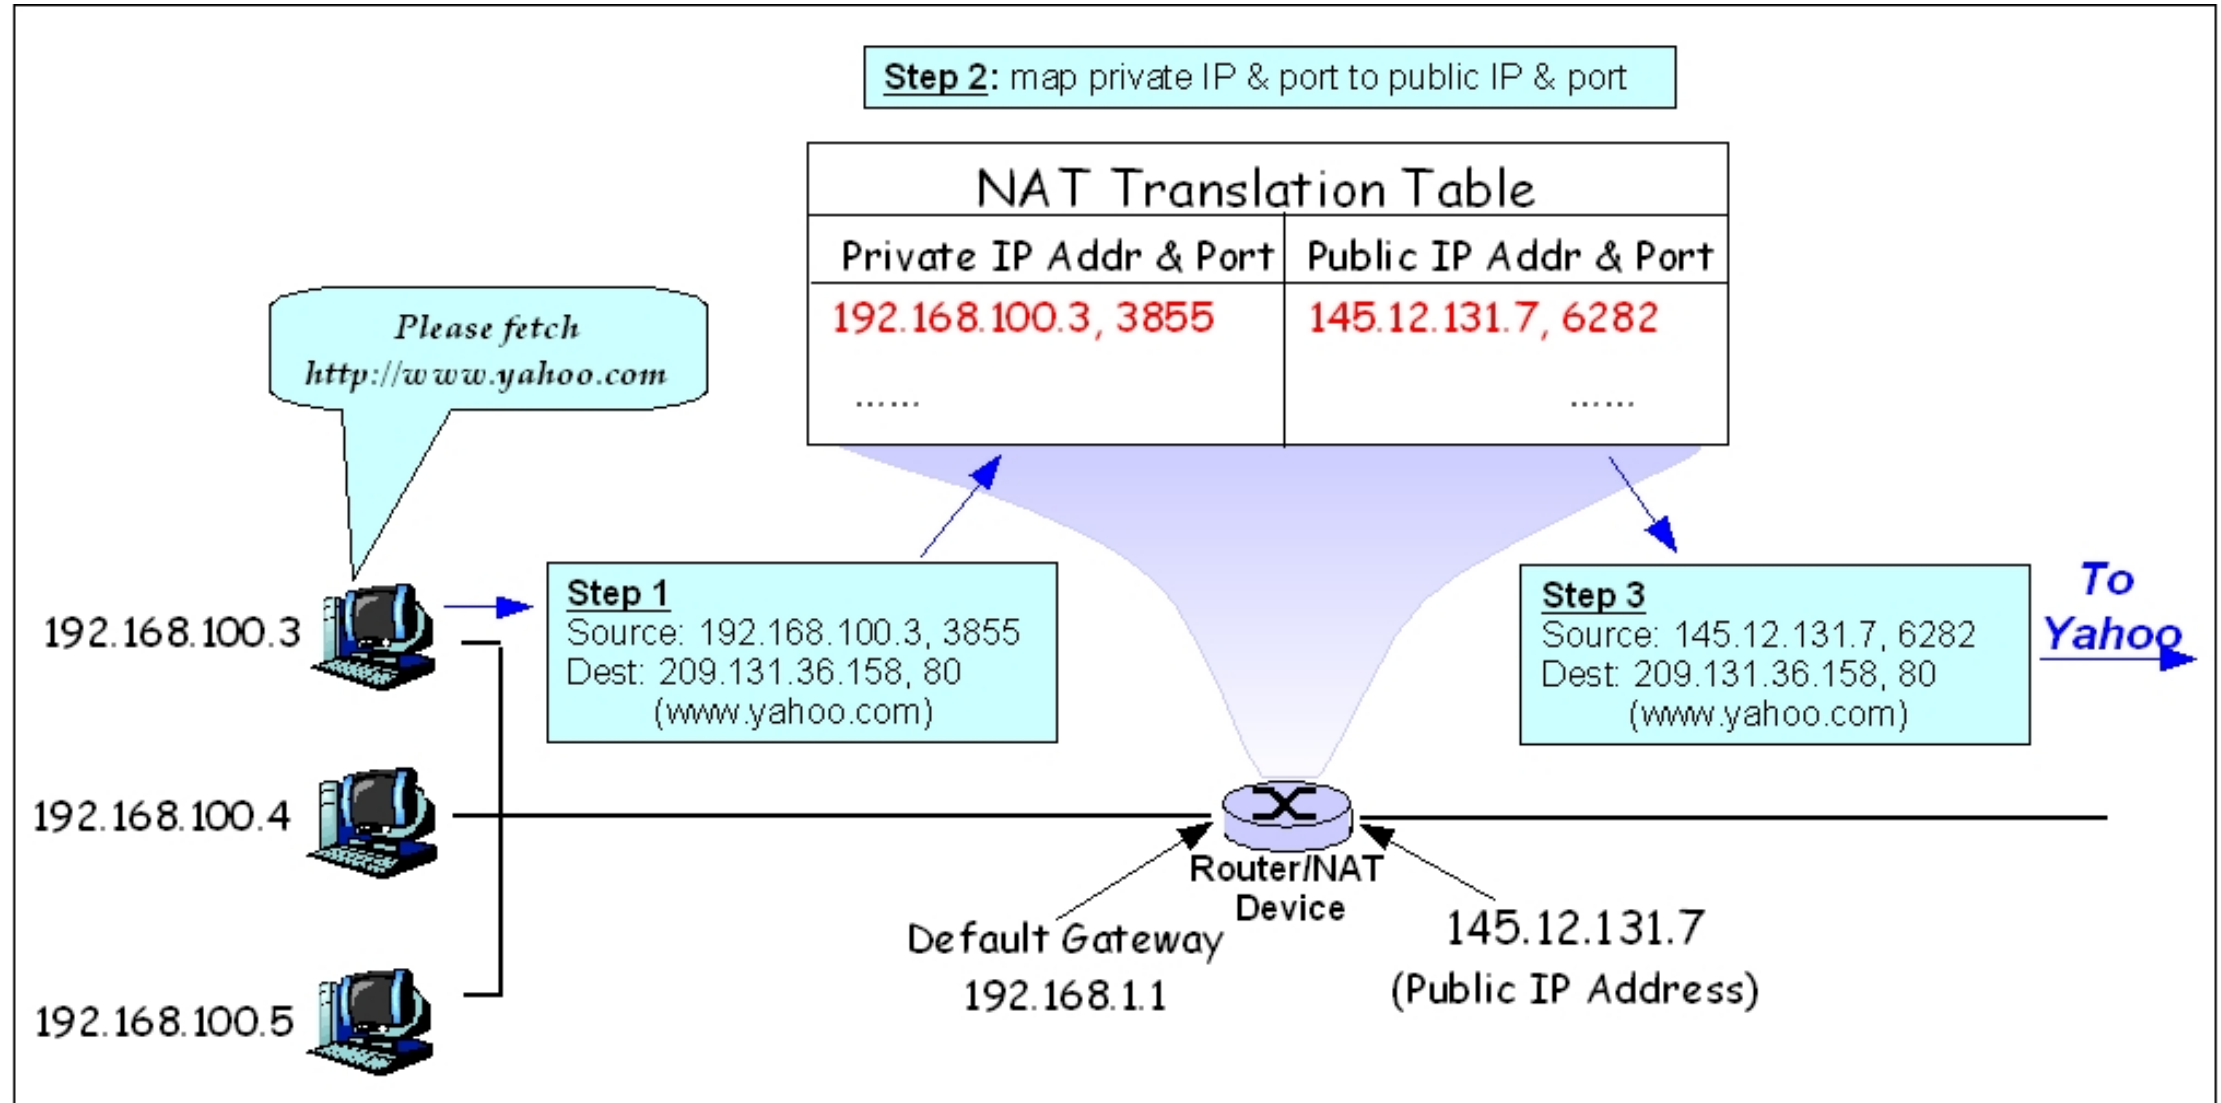
Network

Private Network

- a computer network that uses private IP address space
- anyone may use these addresses without approval from regional or local Internet registries
- private network devices communicate with public internet via [Network Address Translation](#) (NAT)

CIDR notation	IP range
10.0.0.0/8	10.0.0.0 – 10.255.255.255
172.16.0.0/12	172.16.0.0 – 172.31.255.255
192.168.0.0/16	192.168.0.0 – 192.168.255.255

Network



Network

Virtual Network

- Like traditional network but in cloud and dedicated only for you
- Has **address space** specified as CIDR block
- Has 1..N **subnets** (CIDR)
- Can have Route table, security rules/groups, public/private endpoints
- Could be **peered** with other Virtual and on-premise networks

Virtual Network / Virtual Private Cloud (VPC) / Virtual Private Cloud

Network

Load Balancer (LB)

- distributes network traffic
- could be exposed publicly or private-only
- custom rules, auto-scaling, health-checks, and others
- integrates with other cloud-services

Load Balancer / Elastic Load Balancer (ELB) / Cloud Load Balancer

Network

Domain Name System (DNS)


- *translates* names to IP addresses
- can work with cloud-internal resources ...
- ... or with public domain-registrars

Azure DNS / Route 53 / Cloud DNS

Network

Content Delivery Network (CDN)

- Caches files on edge locations

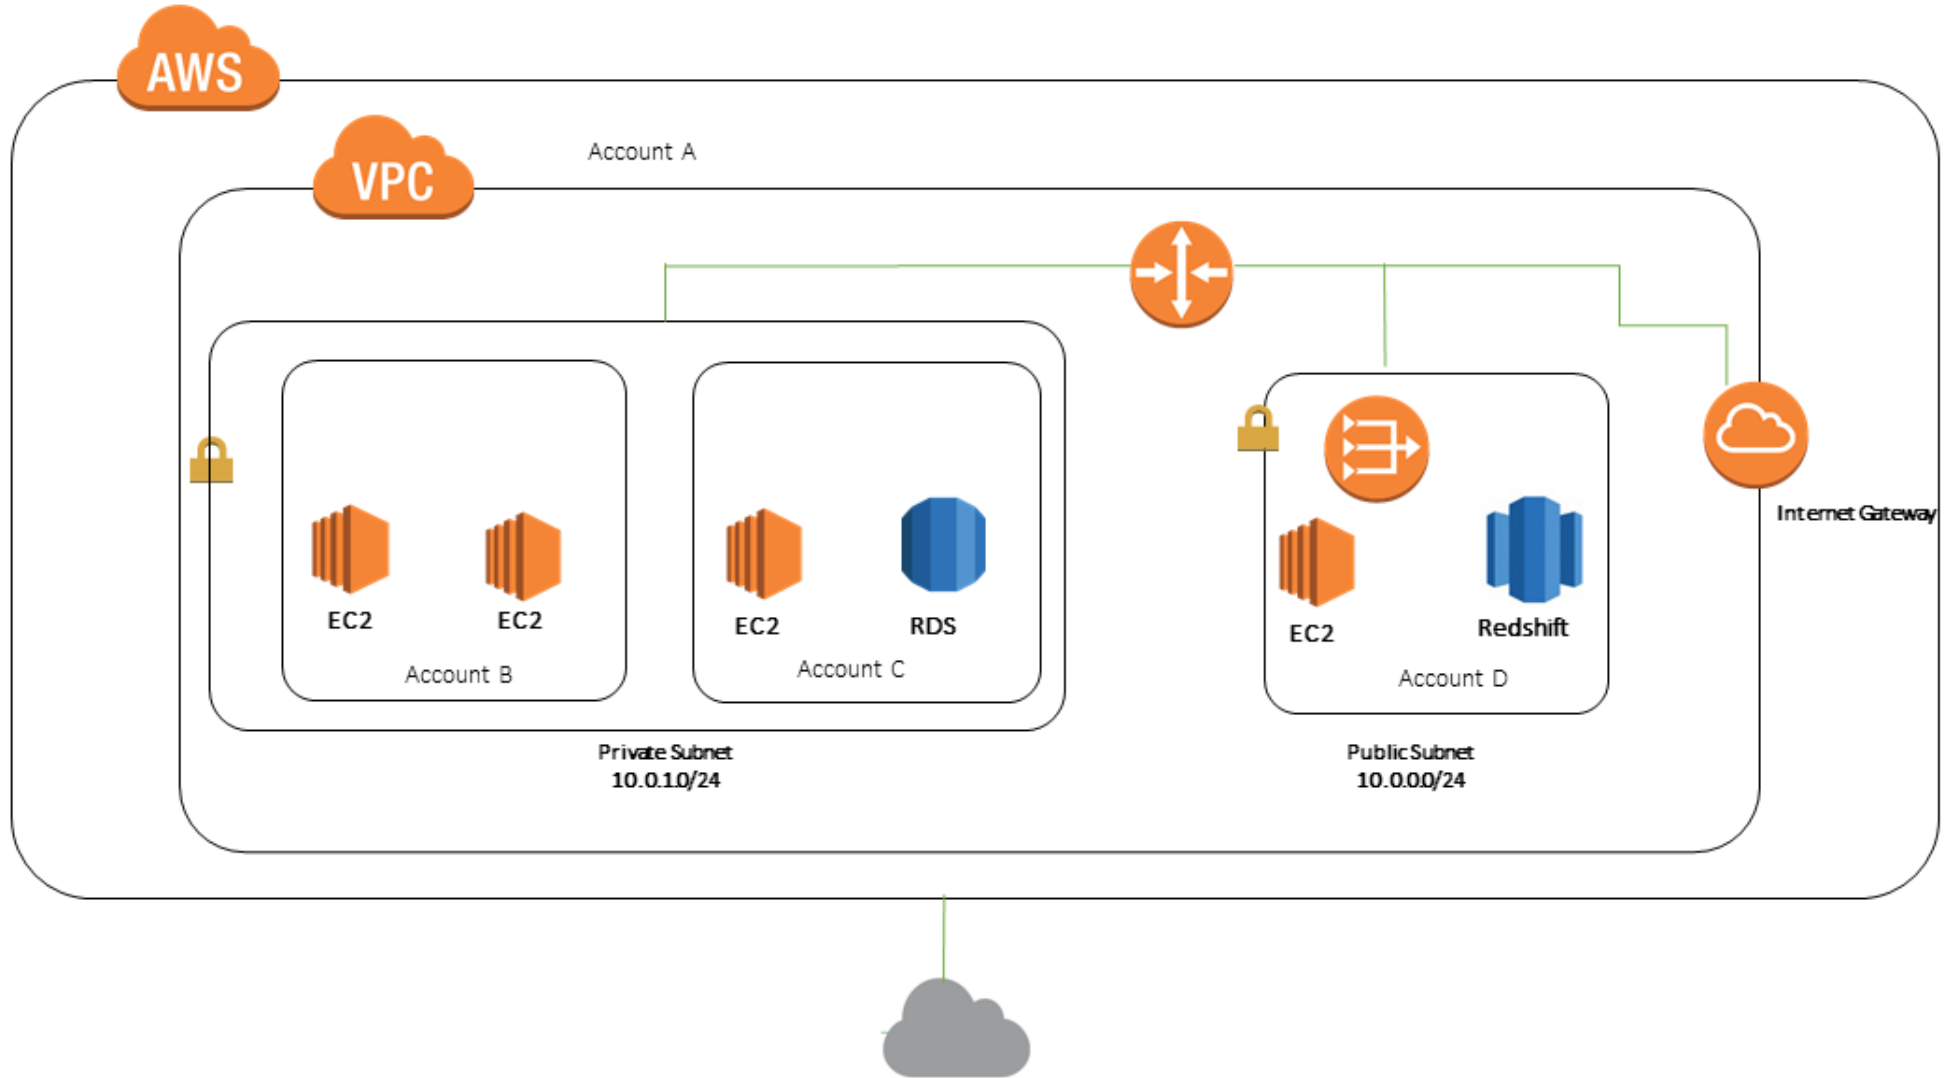


Azure CDN / CloudFront / Cloud CDN

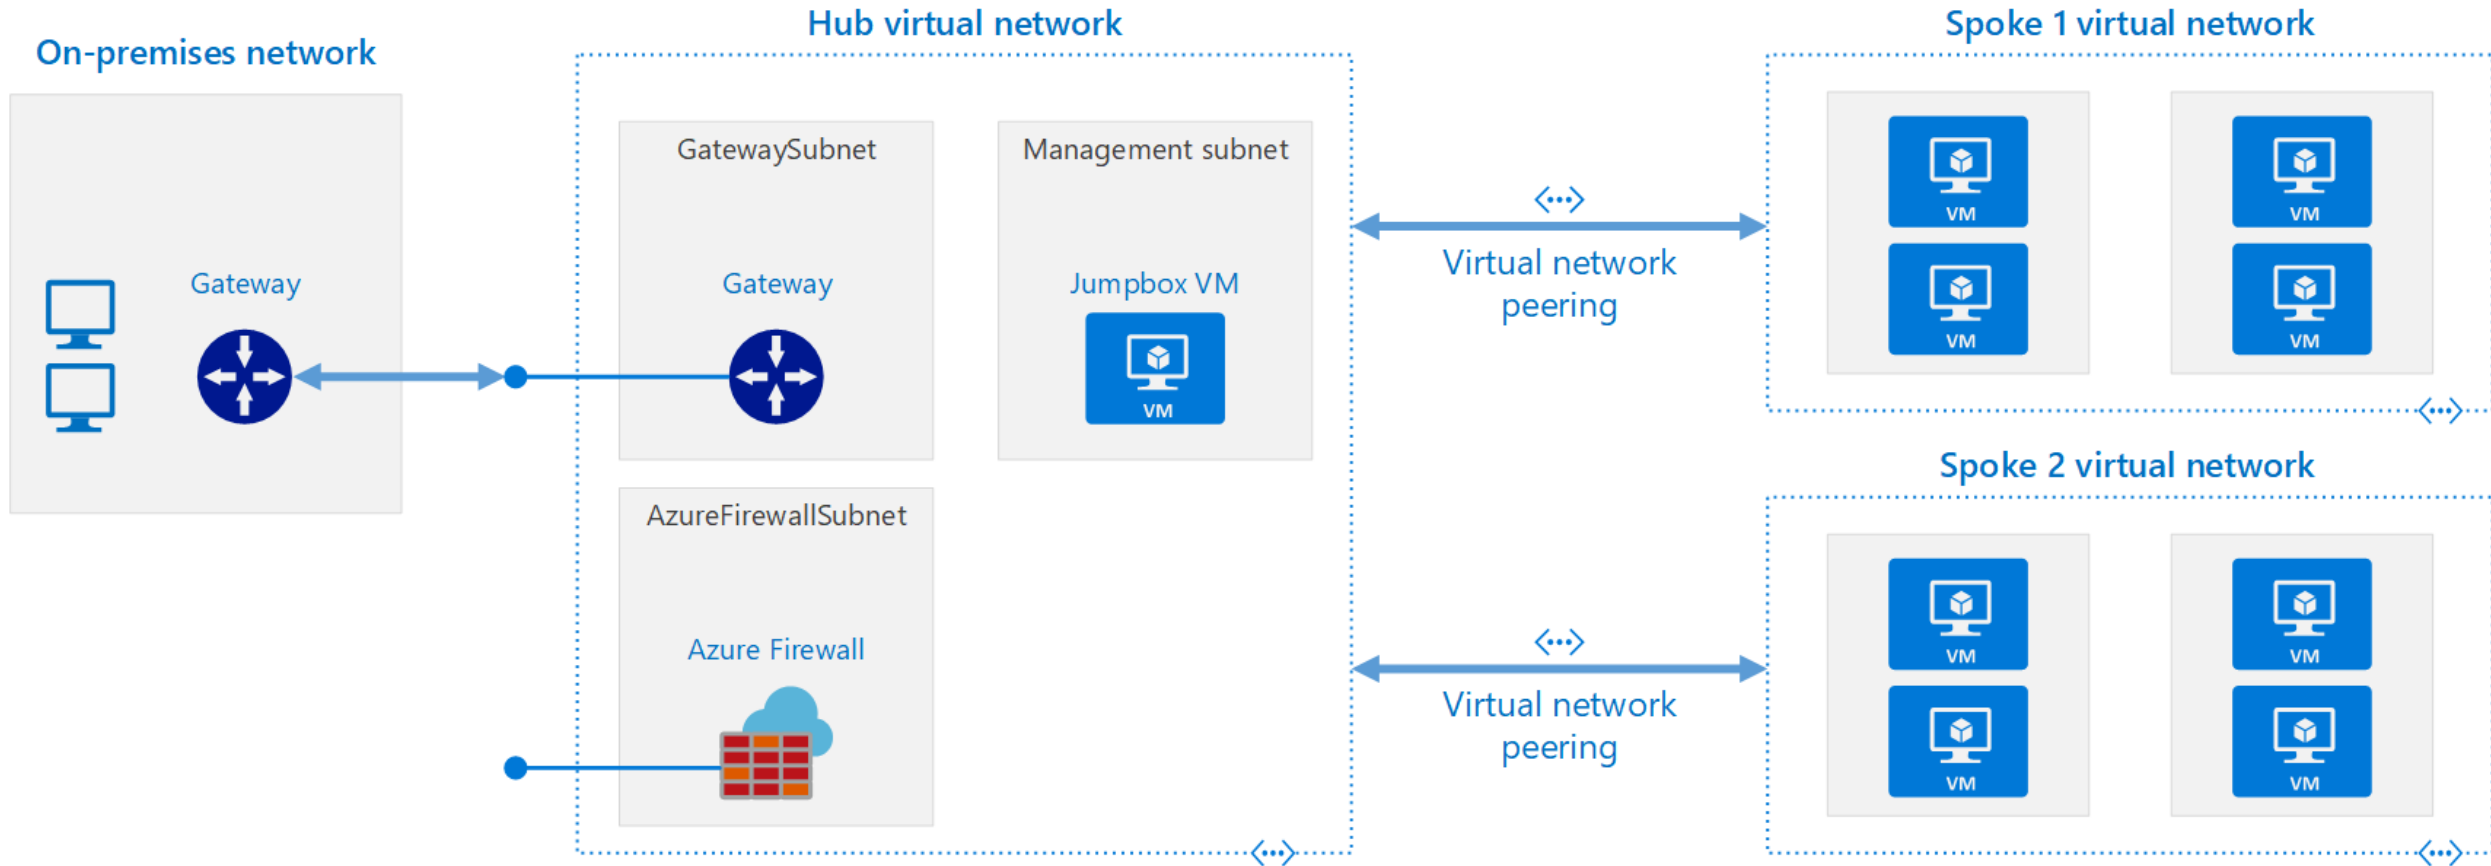
Network

And a LOT of others, for example [AWS Ground Station](#) or [Azure Orbital](#)

Network



Network



Network Summary

- Networking services solve your **connectivity** and **isolation** problems
- Virtual Networks
 - CIDR notation for address space
 - Split on subnets
 - NAT
 - security rules/groups
- Load Balancer
- DNS
- CDN

Databases

Cloud Database is ...

- Place where you store and query data
- Managed by cloud provider

Databases

Cloud providers

- might propose a custom database ([CosmosDB](#), [Aurora](#), [DynamoDB](#), [BigTable](#), [Spanner](#))
- have product, which is compatible with popular solutions: postgresql, mysql, mongo, cassandra, redis, etc
 - **Compatible does not mean actual product behind.** Thus, read the docs

Databases

Cloud Database groups

- Relational
- Non-relational
- In-memory

Databases

Relational DBs in cloud

- Azure:
 - [Azure SQL](#) (MS SQL)
 - [Azure Database for MariaDB/MySQL/PostgreSQL](#)
- AWS:
 - [Aurora](#)
 - [Relational Database Service \[RDS\]](#) (Aurora, psql, mysql, mariadb, oracle, mssql)
- GCP:
 - [Cloud Spanner](#)
 - [Cloud SQL](#) (MySQL, PostgreSQL, and SQL Server)

Databases

Non-relational DBs in cloud

- Azure:
 - [CosmosDB](#) (with different APIs: mongo, cassandra, sql, gremlin)
 - [Storage Account Table](#)
- AWS:
 - [DynamoDB](#)
 - [DocumentDB](#) (mongo), [Keyspaces](#) (cassandra), [Neptune](#) (graph)
- GCP:
 - [Cloud BigTable](#) (more data)
 - [Firestore](#) (less data)

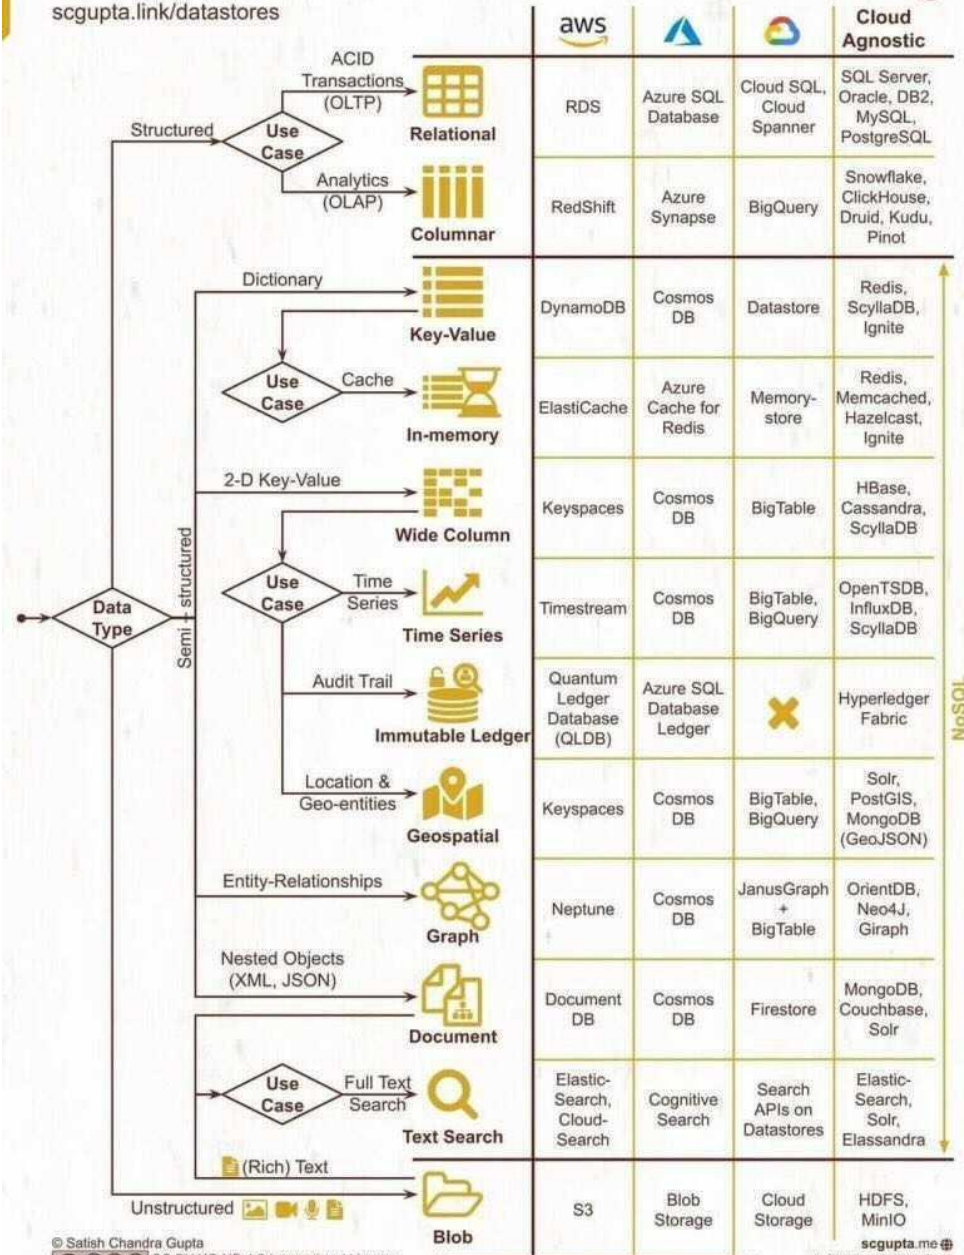
Databases

In-memory DBs in cloud


- Azure:
 - [Azure Cache for Redis](#)
- AWS:
 - [Amazon ElastiCache for Memcached/Redis](#)
- GCP:
 - [Memorystore](#) (Redis/Memcached)

SQL vs. NoSQL: Cheatsheet for AWS, Azure, and Google Cloud

scgupta.link/datastores



Demo



Events and Messaging

Event streams in cloud


- Azure:
 - [Event Hub](#) (compatible with Kafka)
 - [Stream Analytics](#)
- AWS:
 - [Amazon Kinesis](#)
 - [Managed Streaming for Apache Kafka](#)
- GCP:
 - [PubSub](#)
 - [Dataflow](#)


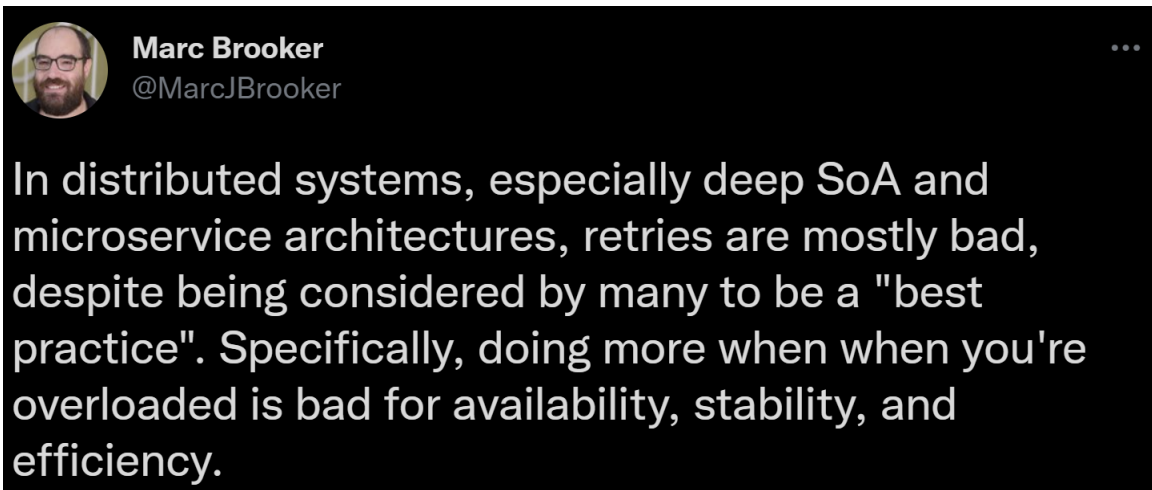
Events and Messaging

Messaging in cloud

- Azure:
 - [Service Bus](#)
 - [Storage Account Queue](#)
- AWS:
 - [Simple Queue Service \(SQS\)](#)
 - [Amazon MQ](#) (ActiveMQ and Rabbit)
- GCP:
 - [PubSub](#)

Demo





<https://twitter.com/MarcJBrooker/status/1489651911640825858>

Additional resources

- (video) [Safe Client Behavior](#)
- (comics) [Networking basics by Julia Evans](#)
- (video) [DynamoDB design-patterns](#) (applies to most “nosql”)

Additional resources (optional)

- (video) [AWS Networking Fundamentals](#)
- (article) [How DNS works](#)
- (article) [What happens on DNS update](#)
- (article) [What is CDN](#)
- (course) Google's [Introduction to Computer Networking](#)

Additional resources (optional)

- (article) [Online Event Processing](#)
- (safe-client articles): [bulkhead](#), [circuit-breaker](#), [poison-queue](#)
- (article) [AWS Aurora inside](#)
- (article) [DynamoDB internals](#) (medium paywall)
- (article) [AWS in plain English](#)